



SPECTRAL NORMALIZATION FOR GENERATIVE ADVERSARIAL NETWORKS

Takeru Miyato, Toshiki Kataoka, Masanori Koyama, Yuichi Yoshida

ICLR 2018

Presenter: Yao Zhang



Introduction of paper

My project and experiment

My contribution

Results

Demo



Introduction

GAN is a framework to produce a model distribution that mimics a given target distribution, and it consists of a generator(G) that produces the model distribution and a discriminator(D) that distinguishes the model distribution from the target.



Introduction

GAN is a framework to produce a model distribution that mimics a given target distribution, and it consists of a generator(G) that produces the model distribution and a discriminator(D) that distinguishes the model distribution from the target.

D and G play the following two-player minimax game with value function $V(G, D)$:

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))].$$



Introduction

D and G play the following two-player minimax game with value function V (G, D):

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))].$$

D maximize: $\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$

G minimize: $\mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$



Introduction

The machine learning community has pointed out that if discriminator is so good during training, the gradient of generator would vanish.

And some work has done to control the training of discriminator. One way is to control the Lipschitz constant of the discriminator function which is proved to be effective.

$$\text{for any } x, y \quad ||D(x) - D(y)|| \leq K||x - y||$$

D is Lipschitz continuous and K is Lipschitz constant of D



Introduction

Spectral normalization controls the Lipschitz constant of the discriminator by constraining the spectral norm of each layer.

Spectral normalization normalizes the spectral norm of the weight matrix W :

$$\bar{W}_{\text{SN}}(W) := W / \sigma(W).$$

If we normalize weight matrix W of each layer, the function of D is Lipschitz continuous and $K = 1$.



Introduction

Spectral normalization controls the Lipschitz constant of the discriminator by constraining the spectral norm of each layer.

Spectral normalization normalizes the spectral norm of the weight matrix W :

$$\bar{W}_{\text{SN}}(W) := W / \sigma(W).$$

If we normalize weight matrix W of each layer, the function of D is Lipschitz continuous and $K = 1$.

By definition, $\sigma(W)$ is equivalent to the largest singular value of W .



Introduction

Spectral normalization controls the Lipschitz constant of the discriminator by constraining the spectral norm of each layer.

Spectral normalization normalizes the spectral norm of the weight matrix W :

$$\bar{W}_{\text{SN}}(W) := W / \sigma(W).$$

If we normalize weight matrix W of each layer, the function of D is Lipschitz continuous and $K = 1$.

By definition, $\sigma(W)$ is equivalent to the largest singular value of W .

Use power iteration to estimate $\sigma(W)$ instead of using SVD.



My project and experiment

I run my code on colab.



My project and experiment

I run my code on colab.

I run a standard DCGAN for CIFAR-10 which is well-designed and it uses batch normalization.



My project and experiment

I run my code on colab.

I run a standard DCGAN for CIFAR-10 which is well-designed and it uses batch normalization.

I revise the DCGAN to use spectral normalization which is my second model. No batch normalization is in this model and only use spectral normalization.



My project and experiment

I run my code on colab.

I run a standard DCGAN for CIFAR-10 which is well-designed and it uses batch normalization.

I revise the DCGAN to use spectral normalization which is my second model. No batch normalization is in this model and only use spectral normalization.

My third model is a DCGAN without using batch normalization or spectral normalization.



My project and experiment

I run my code on colab.

I run a standard DCGAN for CIFAR-10 which is well-designed and it uses batch normalization.

I revise the DCGAN to use spectral normalization which is my second model. No batch normalization is in this model and only use spectral normalization.

My third model is a DCGAN without using batch normalization or spectral normalization.

I use FID score as performance metric.



My project and experiment

I run my code on colab.

I run a standard DCGAN for CIFAR-10 which is well-designed and it uses batch normalization.

I revise the DCGAN to use spectral normalization which is my second model. No batch normalization is in this model and only use spectral normalization.

My third model is a DCGAN without using batch normalization or spectral normalization.

I use FID score as performance metric.

The training time is about one and a half hours for one model.



My project and experiment

Generator of DCGAN

Layer (type)	Output Shape	Param #
ConvTranspose2d-1	[-1, 512, 4, 4]	819,200
BatchNorm2d-2	[-1, 512, 4, 4]	1,024
ReLU-3	[-1, 512, 4, 4]	0
ConvTranspose2d-4	[-1, 256, 8, 8]	2,097,152
BatchNorm2d-5	[-1, 256, 8, 8]	512
ReLU-6	[-1, 256, 8, 8]	0
ConvTranspose2d-7	[-1, 128, 16, 16]	524,288
BatchNorm2d-8	[-1, 128, 16, 16]	256
ReLU-9	[-1, 128, 16, 16]	0
ConvTranspose2d-10	[-1, 64, 32, 32]	131,072
BatchNorm2d-11	[-1, 64, 32, 32]	128
ReLU-12	[-1, 64, 32, 32]	0
ConvTranspose2d-13	[-1, 3, 64, 64]	3,072
Tanh-14	[-1, 3, 64, 64]	0
Total params: 3,576,704		
Trainable params: 3,576,704		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 3.00		
Params size (MB): 13.64		
Estimated Total Size (MB): 16.64		



My project and experiment

Discriminator of DCGAN

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 32, 32]	3,072
LeakyReLU-2	[-1, 64, 32, 32]	0
Conv2d-3	[-1, 128, 16, 16]	131,072
BatchNorm2d-4	[-1, 128, 16, 16]	256
LeakyReLU-5	[-1, 128, 16, 16]	0
Conv2d-6	[-1, 256, 8, 8]	524,288
BatchNorm2d-7	[-1, 256, 8, 8]	512
LeakyReLU-8	[-1, 256, 8, 8]	0
Conv2d-9	[-1, 512, 4, 4]	2,097,152
BatchNorm2d-10	[-1, 512, 4, 4]	1,024
LeakyReLU-11	[-1, 512, 4, 4]	0
Conv2d-12	[-1, 1, 1, 1]	8,192
Sigmoid-13	[-1, 1, 1, 1]	0

Total params: 2,765,568

Trainable params: 2,765,568

Non-trainable params: 0

Input size (MB): 0.05

Forward/backward pass size (MB): 2.31

Params size (MB): 10.55

Estimated Total Size (MB): 12.91

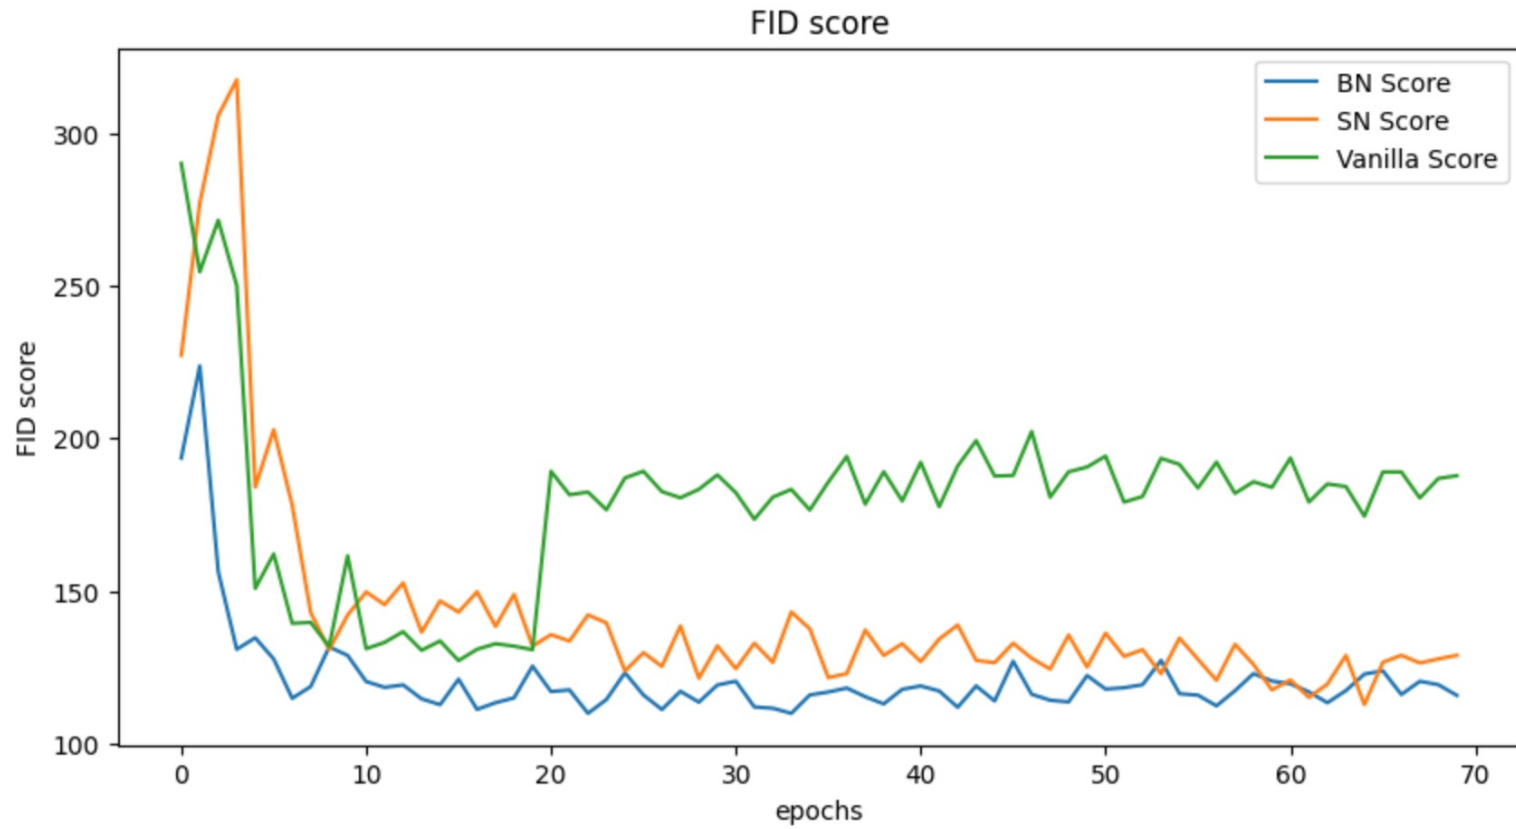


My contribution

The architecture of deep neural networks is different from the architecture used in the paper which can prove the generalization and effectiveness of spectral normalization.




Results





Results


Fake Images





Results


Fake Images with Spectral Normalization





Results

Fake Images without Batch Normalization or Spectral Normalization





Demo

Colab:

https://colab.research.google.com/drive/1_Ktc3XsLH2RpSahY1ly5Sng0EptNQ4VJ?usp=sharing

Github:

https://github.com/yzhan480/CS536Final_Project/